

K. J. Kevin Feng

kjfeng@uw.edu

kjfeng.me

Research Interests

Human-computer interaction; social and collaborative computing; AI and society.

Education

- 2021 – **University of Washington**, Seattle, WA
Ph.D. Human Centered Design & Engineering
Advisors: David W. McDonald (HCDE), Amy X. Zhang (CSE)
- 2021 – 2023 M.S. Human Centered Design & Engineering
- 2017 – 2021 **Princeton University**, Princeton, NJ
B.S.E. Computer Science, Magna cum laude
Thesis: *Lowering the Barrier for Web Advertisement Research at Scale*
Advisor: Arvind Narayanan
Minors: technology & society, visual arts
- 01 – 05 2020 **National University of Singapore**, Singapore
Study Abroad, Computer Science

Publications

* denotes equal contribution. Acceptance rates for venues, when available, are in [square brackets].

REFEREED CONFERENCE AND JOURNAL PAPERS

- [c.9] **K. J. Kevin Feng**, Q. Vera Liao, Ziang Xiao, Jennifer Wortman Vaughan, Amy X. Zhang, David W. McDonald. Canvil: Designerly Adaptation for LLM-Powered User Experiences. *ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025. [25.1%].
- [c.8] Inyoung Cheong, King Xia, **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Amy X. Zhang. (A)I Am Not a Lawyer, But...: Engaging Legal Experts towards Responsible LLM Policies for Legal Advice. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*. Rio de Janeiro, Brazil. June 2024. [22.9%].
- [c.7] **K. J. Kevin Feng**, Xander Koo, Lawrence Tan, Amy S. Bruckman*, David W. McDonald*, Amy X. Zhang*. Mapping the Design Space of Teachable Social Media Feed Experiences. *ACM Conference on Human Factors in Computing Systems (CHI '24)*. Honolulu, HI. May 2024. [26.3%].

- [c.6] **K. J. Kevin Feng**, Nick Ritchie, Pia Blumenthal, Andy Parsons, Amy X. Zhang. Examining the Impact of Provenance-Enabled Media on Trust and Accuracy Perceptions. *Proceedings of the ACM on Human-Computer Interaction (CSCW '23)*. Minneapolis, MN. October 2023. [32.7%].
- [c.5] **K. J. Kevin Feng**, Maxwell James Coppock, David W. McDonald. How Do UX Practitioners Communicate AI as a Design Material? Artifacts, Conceptions, and Propositions. *ACM Conference on Designing Interactive Systems (DIS '23)*. Pittsburgh, PA. July 2023. [24.0%].
- [c.4] **K. J. Kevin Feng**, Tony W. Li, Amy X. Zhang. Understanding Collaborative Practices and Tools of Professional UX Practitioners in Software Organizations. *ACM Conference on Human Factors in Computing Systems (CHI '23)*. Hamburg, Germany. April 2023. [28.4%].
- [c.3] **K. J. Kevin Feng**, David W. McDonald. Addressing UX Practitioners' Challenges in Designing ML Applications: an Interactive Machine Learning Approach. *ACM Conference on Intelligent User Interfaces (IUI '23)*. Sydney, Australia. March 2023. [24.1%].
- [c.2] **K. J. Kevin Feng**, Kevin Song, Kejing Li, Oishee Chakrabarti, Marshini Chetty. Investigating How University Students in the United States Encounter and Deal With Misinformation in Private WhatsApp Chats During COVID-19. *Eighteenth Symposium on Usable Privacy and Security (SOUPS '22)*. Boston, MA. August 2022. [27.8%].
- [c.1] Ruoxi Shang, **K. J. Kevin Feng**, Chirag Shah. Why Am I Not Seeing It? Understanding Lay Users' Needs in Counterfactual Explanations for Everyday Recommendations. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. Seoul, South Korea. June 2022. [25.2%].

WORKSHOP PAPERS AND POSTERS

- [w.5] Andrew Konya, Aviv Ovadya, **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Lisa Schirch, Colin Irwin, Amy X. Zhang. Chain of Alignment: Integrating Public Will with Expert Intelligence for Language Model Alignment. *Pluralistic Alignment workshop @ NeurIPS '24*. Vancouver, Canada. December 2024.
- [w.4] **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Inyoung Cheong, King Xia, Amy X. Zhang. Case Repositories: Towards Case-Based Reasoning for AI Alignment. *AI Meets Moral Philosophy and Moral Psychology (MP2) workshop @ NeurIPS '23*. New Orleans, LA. December 2023.
- [w.3] **K. J. Kevin Feng**, David W. McDonald, Amy X. Zhang. Teachable Agents for End-User Empowerment in Personalized Feed Curation. *Knight First Amendment Institute symposium: Optimizing for What? Algorithmic Amplification and Society*. New York City, NY. April 2023. [18.9%].
- [w.2] **K. J. Kevin Feng***, Alice Gao*, Johanna S. Karras*. Towards Semantically Aware Word Cloud Shape Generation. *Adjunct Proceedings of the ACM Symposium on User Interface Software and Technology (UIST '22)*. Bend, OR. October 2022.

- [w.1] **K. J. Kevin Feng**, Amy X. Zhang. From Handoffs to Co-Creation: Deepening Collaboration between Designers, Developers, and Data Science Workers in UX Design. *InContext Workshop @ CHI '22*. New Orleans, LA. May 2022.

THESES AND NON-REFEREED PAPERS

- [n.1] **K. J. Kevin Feng**. Lowering the Barrier for Web Advertisement Research at Scale. *Princeton University Department of Computer Science*. Princeton, NJ. April 2021.

PREPRINTS

K. J. Kevin Feng, Kevin Pu, Matt Latzke, Tal August, Pao Siangliulue, Jonathan Bragg, Daniel S. Weld, Amy X. Zhang, Joseph Chee Chang. Cocoa: Co-Planning and Co-Execution with AI Agents.

K. J. Kevin Feng, Inyoung Cheong, Quan Ze (Jim) Chen, Amy X. Zhang. Policy Prototyping for LLMs: Pluralistic Alignment via Interactive and Collaborative Policymaking.

Research Experience

- 06 – 08 2024 Research Intern, **Semantic Scholar, Allen Institute for AI**, Seattle, WA
Host: Joseph Chee Chang
Interactive LLM agents in collaborative documents.
- 06 – 09 2023 Research Intern, **Fairness, Accountability, Transparency, and Ethics (FATE) Group, Microsoft Research**, Montréal, Canada
Host: Q. Vera Liao
LLMs as UX design material [c.9].
- 01 2022 – Graduate Research Assistant, **Paul G. Allen School of Computer Science & Engineering, University of Washington**, Seattle, WA
Advisor: Amy X. Zhang
Human-centered tools and methods for AI safety + alignment [c.8, w.4, w.5]; social media feed curation via machine teaching [c.7, w.3].
- 09 2021 – Graduate Research Assistant, **Human Centered Design & Engineering, University of Washington**, Seattle, WA
Advisor: David W. McDonald
Interactive machine learning for UX practitioners to design AI-powered applications [c.3, c.5, c.9].
- 08 2020 – Visiting Undergraduate Researcher, **University of Chicago AIRLab**, Remote
07 2021 *Host: Marshini Chetty*
COVID-19 misinformation in private WhatsApp chats [c.2].

06 2020 – Undergraduate Researcher, **Center for Information Technology Policy**, Princeton, NJ
 06 2021 *Advisor: Arvind Narayanan*
 AdOculus: a search platform and dataset for analyzing ads on the web, aided by computer vision [n.1].

Awards & Honours

2024 University of Washington nominee for Google PhD Fellowship
1 of 4 selected from UW Graduate School.
 2023 OpenAI Democratic Inputs to AI grant
1 of 10 teams selected from 860+ worldwide.
 2022 UW Herbold Fellowship
1 of 7 selected from UW College of Engineering.
 2020 Princeton Council for Science and Technology Independent Project Award
 2020 Princeton Office of Undergraduate Research Summer Research Award
 2019 IDEO CoLab Fellowship Finalist
 2018 Keller Center for Innovation in Engineering Education Summer Fellow

Invited Talks & Panels

2025 **Cocoa: Co-Planning and Co-Execution with AI Agents**
 Guest lecture, UW INFO 498: Understanding and Navigating Cognitive and Behavioral Change in the AI Era
 2024 **Mapping the Design Space of Teachable Social Media Feed Experiences**
 Princeton HCI
Canvil: Supporting Designerly Model Adaptation in LLM-Powered User Experiences
 Guest lecture, UW HCDE 598: Human-Centeredness in Generative AI
Case Repositories: Towards Case-Based Reasoning for AI Alignment
 OpenAI Forum: Collective Alignment
 2023 **Teachable Agents for End-User Empowerment in Personalized Feed Curation**
 Machine Intelligence and Normative Theory (MINT) Lab, Australian National University
Canvil: Supporting Designerly Model Adaptation in LLM-Powered User Experiences
 UW HCDE Academic Research Seminar
 Microsoft Research Montréal
 Microsoft Research FATE Group
 Microsoft Research HCI Seminar
[Panel] Empirical look at user behavior
Optimizing for What? Algorithmic Amplification and Society, Knight First Amendment Institute
 2022 **User Perceptions of Media in Provenance-Enabled Social Media Feeds**
 Coalition for Content Provenance and Authenticity (C2PA) Technical Working Group
 2021 **Lowering the Barrier for Web Advertisement Research at Scale**
 Princeton Research Day
 Princeton CITP Works in Progress Reading Group

2020 **User Interface Design, Prototyping, and Testing**
Princeton ACM

2019 **Introduction to User Interface Design**
HackPrinceton 2019

2018 **Introduction to User Interface Design**
HackPrinceton 2018

Service

2021 – **Program Committee:** FAccT 2024–2025
Reviewing: CHI 2023–2025*, NeurIPS Pluralistic Alignment Workshop 2024, CHI LBW 2024*, CSCW 2023–2024*, DIS 2023*, UIST 2022
** recognition for outstanding review(s).*

Student Volunteer: CHI 2023, IUI 2022, UIST 2022

Departmental Service: HCDE PhD Visit Days session leader, HCDE PhD orientation student panelist, HCDE PhD application reviewer

2020 – 2021 **Organizer**, Princeton intersession course on ethical technology

2019 – 2021 **First-year engineering student advisor**, Princeton School of Engineering and Applied Science

2018 – 2021 **Founder**, Rehack (reverse hackathon for development of more socially responsible technologies)

Mentoring

2025– **Jessica Do; Nancy Qu; Aditi Rana, Caleb West**, UW CSE BS, working on analyzing custom feeds on Bluesky.

2023–2024 **Lawrence Tan**, UW CSE BS, co-authored [c.7].

2022–2023 **Tony Li**, UW HCDE MS, co-authored [c.4].
Afterwards: PhD student at UCSD CSE.

2022–2023 **Maxwell Coppock**, UW HCDE BS, co-authored [c.5].
Afterwards: MS student at UW MHCI+D.

2022 **Amodini Khade**, UW HCDE MS.

2022 **Katrina Vergara**, UW HCDE MS.
Afterwards: Product Manager at Microsoft.

Industry Work Experience

06 – 09 2022 Program Manager Intern, **Responsible AI, Azure Machine Learning, Microsoft**, Bellevue, WA
Conducted research on the role of machine learning scorecards in facilitating multi-stakeholder conversations about model deployment.

06 – 08 2021 Program Manager Intern, **Data Platforms + Growth, Microsoft**, Remote
Launched UI for data analysts to better organize and take action on their data.

06 – 08 2019 Product Management Intern, **The Muse**, New York, NY

Built and launched new job search recommendation feature with a team of 2 engineers and a designer.

o6 – o8 2018

Product Design Intern, **Solomoto**, Tel Aviv, Israel

Data dashboard design for small business management SaaS platform.

Skills

Technical: Python, JavaScript, Java, Unity, C, OCaml, Go, SQL, \LaTeX .

Other: Figma, Adobe Creative Suite, qualitative and quantitative research methods, laser cutting, art exhibition installation, bookbinding, photography.