

K. J. Kevin Feng

kjfeng@uw.edu

kjfeng.me

Education

- 2021 – 2026 **University of Washington**, Seattle, WA
Ph.D. Human Centered Design & Engineering
Dissertation: *Designing for Interactive Systems Powered by Generally Capable AI*
Advisors: David W. McDonald (HCDE), Amy X. Zhang (CSE)
- 2021 – 2023 M.S. Human Centered Design & Engineering
- 2017 – 2021 **Princeton University**, Princeton, NJ
B.S.E. Computer Science, Magna cum laude
Thesis: *Lowering the Barrier for Web Advertisement Research at Scale*
Advisor: Arvind Narayanan
Minors: technology & society, visual arts
- 01 – 05 2020 **National University of Singapore**, Singapore
Study Abroad, Computer Science

Publications

* denotes equal contribution. Acceptance rates for venues, when available, are in [square brackets].

REFEREED CONFERENCE AND JOURNAL PAPERS

- [c.15] Leon Staufer, **K. J. Kevin Feng**, Kevin Wei, Luke Bailey, Yawen Duan, Mick Yang, A. Pinar Ozisik, Stephen Casper, Noam Kolt. The 2025 AI Agent Index: Documenting Technical and Safety Features of Deployed Agentic AI Systems. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '26)*. Montréal, Canada. April 2026. [33.0%].
- [c.14] **K. J. Kevin Feng**, Kevin Pu, Matt Latzke, Tal August, Pao Siangliulue, Jonathan Bragg, Daniel S. Weld, Amy X. Zhang, Joseph Chee Chang. Cocoa: Co-Planning and Co-Execution with AI Agents. *ACM Conference on Human Factors in Computing Systems (CHI '26)*. Barcelona, Spain. April 2026. [25.3%]. **Best Paper** (top 1% of submissions).
- [c.13] **K. J. Kevin Feng**, Tzu-Sheng Kuo, Quan Ze (Jim) Chen, Inyoung Cheong, Kenneth Holstein, Amy X. Zhang. PolicyPad: Collaborative Prototyping of LLM Policies. *ACM Conference on Human Factors in Computing Systems (CHI '26)*. Barcelona, Spain. April 2026. [25.3%].
- [c.12] **K. J. Kevin Feng**, Tzu-Sheng Kuo, Quan Ze (Jim) Chen, Inyoung Cheong, Kenneth Holstein, Amy X. Zhang. Expert-Driven Prototyping of LLM Behavioral Policies: Case Studies in Mental Health

and Law. *2nd Conference by the International Association for Safe & Ethical AI (IASEAI '26)*. Paris, France. February 2026.

- [c.11] Quan Ze (Jim) Chen, **K. J. Kevin Feng**, Chan Young Park, Amy X. Zhang. SPICA: Retrieving Scenarios for Pluralistic In-Context Alignment. *Findings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL '25)*. Vienna, Austria. July 2025.
- [c.10] **K. J. Kevin Feng**, Q. Vera Liao, Ziang Xiao, Jennifer Wortman Vaughan, Amy X. Zhang, David W. McDonald. Canvil: Designerly Adaptation for LLM-Powered User Experiences. *ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025. [24.9%].
- [c.9] Kevin Pu, **K. J. Kevin Feng**, Tovi Grossman, Tom Hope, Bhavana Dalvi Mishra, Matt Latzke, Jonathan Bragg, Joseph Chee Chang, Pao Siangliulue. IdeaSynth: Iterative Research Idea Development Through Evolving and Composing Idea Facets with Literature-Grounded Feedback. *ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025. [24.9%].
- [c.8] Inyoung Cheong, King Xia, **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Amy X. Zhang. (A)I Am Not a Lawyer, But...: Engaging Legal Experts towards Responsible LLM Policies for Legal Advice. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*. Rio de Janeiro, Brazil. June 2024. [22.9%].
- [c.7] **K. J. Kevin Feng**, Xander Koo, Lawrence Tan, Amy S. Bruckman*, David W. McDonald*, Amy X. Zhang*. Mapping the Design Space of Teachable Social Media Feed Experiences. *ACM Conference on Human Factors in Computing Systems (CHI '24)*. Honolulu, HI. May 2024. [26.3%].
- [c.6] **K. J. Kevin Feng**, Nick Ritchie, Pia Blumenthal, Andy Parsons, Amy X. Zhang. Examining the Impact of Provenance-Enabled Media on Trust and Accuracy Perceptions. *Proceedings of the ACM on Human-Computer Interaction (CSCW '23)*. Minneapolis, MN. October 2023. [32.7%].
- [c.5] **K. J. Kevin Feng**, Maxwell James Coppock, David W. McDonald. How Do UX Practitioners Communicate AI as a Design Material? Artifacts, Conceptions, and Propositions. *ACM Conference on Designing Interactive Systems (DIS '23)*. Pittsburgh, PA. July 2023. [24.0%].
- [c.4] **K. J. Kevin Feng**, Tony W. Li, Amy X. Zhang. Understanding Collaborative Practices and Tools of Professional UX Practitioners in Software Organizations. *ACM Conference on Human Factors in Computing Systems (CHI '23)*. Hamburg, Germany. April 2023. [28.4%].
- [c.3] **K. J. Kevin Feng**, David W. McDonald. Addressing UX Practitioners' Challenges in Designing ML Applications: an Interactive Machine Learning Approach. *ACM Conference on Intelligent User Interfaces (IUI '23)*. Sydney, Australia. March 2023. [24.1%].
- [c.2] **K. J. Kevin Feng**, Kevin Song, Kejing Li, Oishee Chakrabarti, Marshini Chetty. Investigating How University Students in the United States Encounter and Deal With Misinformation in Private What-

sApp Chats During COVID-19. *Eighteenth Symposium on Usable Privacy and Security (SOUPS '22)*. Boston, MA. August 2022. [27.8%].

- [c.1] Ruoxi Shang, **K. J. Kevin Feng**, Chirag Shah. Why Am I Not Seeing It? Understanding Lay Users' Needs in Counterfactual Explanations for Everyday Recommendations. *ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. Seoul, South Korea. June 2022. [25.2%].

PREPRINTS

- [p.1] Micah Carroll, Adeline Foote, **K. J. Kevin Feng**, Marcus Williams, Anca Dragan, W. Bradley Knox*, Smitha Milli*. CTRL-Rec: Controlling Recommender Systems With Natural Language. <https://arxiv.org/abs/2510.12742>.

LIGHTLY REFEREED PAPERS AND POSTERS

- [l.8] **K. J. Kevin Feng**, Tae Soo Kim, Rock Yuren Pang, Faria Huq, Tal August, Amy X. Zhang. On the Regulatory Potential of User Interfaces for AI Agent Governance. *Regulatable ML workshop @ NeurIPS '25 (Oral)*. San Diego, CA. December 2025.
- [l.7] **K. J. Kevin Feng**, Inyoung Cheong, Quan Ze (Jim) Chen, Amy X. Zhang. Policy Prototyping for LLMs: Pluralistic Alignment via Interactive and Collaborative Policymaking. *Bidirectional Human-AI Alignment workshop @ ICLR '25 (Oral)*. Singapore. April 2025.
- [l.6] **K. J. Kevin Feng**, David W. McDonald, Amy X. Zhang. Levels of Autonomy for AI Agents. *Knight First Amendment Institute symposium: AI and Democratic Freedoms*. New York City, NY. April 2025.
- [l.5] Andrew Konya, Aviv Ovadya, **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Lisa Schirch, Colin Irwin, Amy X. Zhang. Chain of Alignment: Integrating Public Will with Expert Intelligence for Language Model Alignment. *Pluralistic Alignment workshop @ NeurIPS '24*. Vancouver, Canada. December 2024.
- [l.4] **K. J. Kevin Feng**, Quan Ze (Jim) Chen, Inyoung Cheong, King Xia, Amy X. Zhang. Case Repositories: Towards Case-Based Reasoning for AI Alignment. *AI Meets Moral Philosophy and Moral Psychology (MP2) workshop @ NeurIPS '23*. New Orleans, LA. December 2023.
- [l.3] **K. J. Kevin Feng**, David W. McDonald, Amy X. Zhang. Teachable Agents for End-User Empowerment in Personalized Feed Curation. *Knight First Amendment Institute symposium: Optimizing for What? Algorithmic Amplification and Society*. New York City, NY. April 2023. [18.9%].
- [l.2] **K. J. Kevin Feng***, Alice Gao*, Johanna S. Karras*. Towards Semantically Aware Word Cloud Shape Generation. *Adjunct Proceedings of the ACM Symposium on User Interface Software and Technology (UIST '22)*. Bend, OR. October 2022.

- [l.i] **K. J. Kevin Feng**, Amy X. Zhang. From Handoffs to Co-Creation: Deepening Collaboration between Designers, Developers, and Data Science Workers in UX Design. *InContext Workshop @ CHI '22*. New Orleans, LA. May 2022.

WORKSHOPS ORGANIZED

- [w.i] **K. J. Kevin Feng***, Rock Yuren Pang*, Tzu-Sheng Kuo, Amy A. Winecoff, Emily Tseng, David Gray Widder, Harini Suresh, Katharina Reinecke, Amy X. Zhang. Sociotechnical AI Governance: Challenges and Opportunities for *Workshop at the ACM Conference on Human Factors in Computing Systems (CHI '25)*. Yokohama, Japan. April 2025.

THESES AND NON-REFEREED PAPERS

- [n.i] **K. J. Kevin Feng**. Lowering the Barrier for Web Advertisement Research at Scale. *Princeton University Department of Computer Science*. Princeton, NJ. April 2021.

Research Experience

- 01 – 04 2026 Student Researcher, **Google Research**, Mountain View, CA
Host: Kornnaphop (Ken) Kawintiranon
Multi-agent debate for multimodal AI safety and alignment.
- 06 – 08 2025 Summer Fellow, **Centre for the Governance of AI**, London, United Kingdom
Collaborators: Risk Management Team (GovAI); Human Influence Team (UK AISI)
Chain-of-thought monitoring for harmful manipulation from advanced AI.
- 06 – 08 2024 Research Intern, **Asta Team, Allen Institute for AI**, Seattle, WA
Host: Joseph Chee Chang
Co-planning and co-execution with AI agents in scientific research [c.14].
- 06 – 09 2023 Research Intern, **Fairness, Accountability, Transparency, and Ethics (FATE) Group, Microsoft Research**, Montréal, Canada
Host: Q. Vera Liao
LLMs as UX design material [c.9].
- 01 2022 – Graduate Research Assistant, **Paul G. Allen School of Computer Science & Engineering, University of Washington**, Seattle, WA
Advisor: Amy X. Zhang
Human-centered tools and methods for AI safety + alignment [c.11–13, c.8, w.4, w.5]; social media feed curation via machine teaching [c.7, w.3].
- 09 2021 – Graduate Research Assistant, **Human Centered Design & Engineering, University of Washington**, Seattle, WA

Advisor: David W. McDonald

Interactive machine learning for UX practitioners to design AI-powered applications [c.3, c.5, c.9].

08 2020 –
07 2021

Visiting Undergraduate Researcher, **University of Chicago AIRLab**, Remote

Host: Marshini Chetty

COVID-19 misinformation in private WhatsApp chats [c.2].

06 2020 –
06 2021

Undergraduate Researcher, **Center for Information Technology Policy**, Princeton, NJ

Advisor: Arvind Narayanan

AdOculus: a search platform and dataset for analyzing ads on the web, aided by computer vision [n.1].

Awards & Honours

2026

Best Paper Award, ACM CHI

2024

University of Washington nominee for Google PhD Fellowship

1 of 4 selected from UW Graduate School.

2023

OpenAI Democratic Inputs to AI grant

1 of 10 teams selected from 860+ worldwide.

2022

UW Herbold Fellowship

1 of 7 selected from UW College of Engineering.

2020

Princeton Council for Science and Technology Independent Project Award

2020

Princeton Office of Undergraduate Research Summer Research Award

2019

IDEO CoLab Fellowship Finalist

2018

Keller Center for Innovation in Engineering Education Summer Fellow

Invited Talks & Panels

2026

Responsible AI via Human-Human, Human-AI, and AI-AI Collaboration

Google DeepMind

Designing for Human-AI Collaboration Across Levels of Autonomy

Apple AI/ML

Expert-Informed LLM Policies for High-Stakes Alignment

Snorkel AI

Handshake AI Research

Human-Agent Collaboration Across Levels of Autonomy

Google

2025

Human-Agent Collaboration Across Levels of Autonomy

Databricks

Human-Agent Collaboration: a Grand Challenge in the Agentic AI Era

UW CSE Industry Affiliates Showcase

Chain-of-Thought Monitoring for Harmful Manipulation from Advanced AI

UK AISI Societal Impacts Team

Designing for Interaction with Advanced AI

The Design Lab, UCSD

[Panel] AI Agents' Democratic and Economic Impacts

Artificial Intelligence and Democratic Freedoms, Knight First Amendment Institute

Cocoa: Co-Planning and Co-Execution with AI Agents

Guest lecture, UW INFO 498: Understanding and Navigating Cognitive and Behavioral Change in the AI Era

2024

Mapping the Design Space of Teachable Social Media Feed Experiences

Princeton HCI

Canvil: Supporting Designerly Model Adaptation in LLM-Powered User Experiences

Guest lecture, UW HCDE 598: Human-Centeredness in Generative AI

Case Repositories: Towards Case-Based Reasoning for AI Alignment

OpenAI Forum: Collective Alignment

2023

Teachable Agents for End-User Empowerment in Personalized Feed Curation

Machine Intelligence and Normative Theory (MINT) Lab, Australian National University

Canvil: Supporting Designerly Model Adaptation in LLM-Powered User Experiences

UW HCDE Academic Research Seminar

Microsoft Research Montréal

Microsoft Research FATE Group

Microsoft Research HCI Seminar

[Panel] Empirical look at user behavior

Optimizing for What? Algorithmic Amplification and Society, Knight First Amendment Institute

2022

User Perceptions of Media in Provenance-Enabled Social Media Feeds

Coalition for Content Provenance and Authenticity (C2PA) Technical Working Group

2021

Lowering the Barrier for Web Advertisement Research at Scale

Princeton Research Day

Princeton CITP Works in Progress Reading Group

2020

User Interface Design, Prototyping, and Testing

Princeton ACM

2019

Introduction to User Interface Design

HackPrinceton 2019

2018

Introduction to User Interface Design

HackPrinceton 2018

Service

2021 –

Program Committee: FAccT 2024–2025

Reviewing: CHI 2023–2025*, NeurIPS Pluralistic Alignment Workshop 2024, CHI LBW 2024*, CSCW 2023–2024*, DIS 2023*, UIST 2022

* *recognition for outstanding review(s)*.

Student Volunteer: CHI 2023, IUI 2022, UIST 2022

Departmental Service: HCDE PhD Visit Days session leader, HCDE PhD orientation student panelist, HCDE PhD application reviewer

2020 – 2021 **Organizer**, Princeton intersession course on ethical technology
 2019 – 2021 **First-year engineering student advisor**, Princeton School of Engineering and Applied Science
 2018 – 2021 **Founder**, Rehack (reverse hackathon for development of more socially responsible technologies)

Mentoring

2025– **Jessica Do; Nancy Qu; Aditi Rana, Caleb West**, UW CSE BS, working on analyzing custom feeds on Bluesky.
 2023–2024 **Lawrence Tan**, UW CSE BS, co-authored [c.7].
 2022–2023 **Tony Li**, UW HCDE MS, co-authored [c.4].
Afterwards: PhD student at UCSD CSE.
 2022–2023 **Maxwell Coppock**, UW HCDE BS, co-authored [c.5].
Afterwards: MS student at UW MHCI+D.
 2022 **Amodini Khade**, UW HCDE MS.
 2022 **Katrina Vergara**, UW HCDE MS.
Afterwards: Product Manager at Microsoft.

Industry Work Experience

06 – 09 2022 Program Manager Intern, **Responsible AI, Azure Machine Learning, Microsoft**, Bellevue, WA
 Conducted research on the role of machine learning scorecards in facilitating multi-stakeholder conversations about model deployment.
 06 – 08 2021 Program Manager Intern, **Data Platforms + Growth, Microsoft**, Remote
 Launched UI for data analysts to better organize and take action on their data.
 06 – 08 2019 Product Management Intern, **The Muse**, New York, NY
 Built and launched new job search recommendation feature with a team of 2 engineers and a designer.
 06 – 08 2018 Product Design Intern, **Solomoto**, Tel Aviv, Israel
 Data dashboard design for small business management SaaS platform.

Skills

Technical: Python, JavaScript, Java, Unity, C, OCaml, Go, SQL, $\text{\LaTeX}$.

Other: Figma, Adobe Creative Suite, qualitative and quantitative research methods, laser cutting, art exhibition installation, bookbinding, photography.